

B-test: A Non-parametric, Low Variance Kernel Two-sample Test

Authored by:

Arthur Gretton
Matthew Blaschko
Wojciech Zaremba

Abstract

We propose a family of maximum mean discrepancy (MMD) kernel two-sample tests that have low sample complexity and are consistent. The test has a hyperparameter that allows one to control the tradeoff between sample complexity and computational time. Our family of tests, which we denote as B-tests, is both computationally and statistically efficient, combining favorable properties of previously proposed MMD two-sample tests. It does so by better leveraging samples to produce low variance estimates in the finite sample case, while avoiding a quadratic number of kernel evaluations and complex null-hypothesis approximation as would be required by tests relying on one sample U-statistics. The B-test uses a smaller than quadratic number of kernel evaluations and avoids completely the computational burden of complex null-hypothesis approximation while maintaining consistency and probabilistically conservative thresholds on Type I error. Finally, recent results of combining multiple kernels transfer seamlessly to our hypothesis test, allowing a further increase in discriminative power and decrease in sample complexity.

1 Paper Body

Given two samples $\{x_i\}_{i=1}^n$ where $x_i \sim P$ i.i.d., and $\{y_i\}_{i=1}^n$, where $y_i \sim Q$ i.i.d, the two sample problem consists in testing whether to accept or reject the null hypothesis H_0 that $P = Q$, vs the alternative hypothesis H_A that P and Q are different. This problem has recently been addressed using measures of similarity computed in a reproducing kernel Hilbert space (RKHS), which apply in very general settings where P and Q might be distributions over high dimensional data or structured objects. Kernel test statistics include the maximum mean discrepancy [10, 6] (of which the energy distance is an example [18, 2, 22]), which is the distance between expected features of P and Q in the RKHS; the kernel Fisher discriminant [12], which is the distance between expected feature maps normalized by the feature space covariance; and density ratio estimates

[24]. When used in testing, it is necessary to determine whether the empirical estimate of the relevant similarity measure is sufficiently large as to give the hypothesis $P = Q$ low probability; i.e., below a user-defined threshold α , denoted the test level. The test power denotes the probability of correctly rejecting the null hypothesis, given that $P \neq Q$. The minimum variance unbiased estimator MMDu of the maximum mean discrepancy, on the basis of n samples observed from each of P and Q , is a U-statistic, costing $O(n^2)$ to compute. Unfortunately, this statistic is degenerate under the null hypothesis H_0 that $P = Q$, and its asymptotic distribution takes the form of an infinite weighted sum of independent χ^2 variables (it is asymptotically Gaussian under the alternative hypothesis H_A that $P \neq Q$). Two methods for empirically estimating the null distribution in a consistent way have been proposed: the bootstrap [10], and a method requiring an eigendecomposition of the kernel matrices computed on the merged samples from P and Q [7]. Unfortunately, both procedures are computationally demanding: the former costs $O(n^2)$, with a large constant (the MMD must be computed repeatedly over random assignments of the pooled data); the latter costs $O(n^3)$, but with a smaller constant, hence can in practice be 1

faster than the bootstrap. Another approach is to approximate the null distribution by a member of a simpler parametric family (for instance, a Pearson curve approximation), however this has no consistency guarantees. More recently, an $O(n)$ unbiased estimate MMDl of the maximum mean discrepancy has been proposed [10, Section 6], which is simply a running average over independent pairs of samples from P and Q . While this has much greater variance than the U-statistic, it also has a simpler null distribution: being an average over i.i.d. terms, the central limit theorem gives an asymptotically Normal distribution, under both H_0 and H_A . It is shown in [9] that this simple asymptotic distribution makes it easy to optimize the Hodges and Lehmann asymptotic relative efficiency [19] over the family of kernels that define the statistic: in other words, to choose the kernel which gives the lowest Type II error (probability of wrongly accepting H_0) for a given Type I error (probability of wrongly rejecting H_0). Kernel selection for the U-statistic is a much harder question due to the complex form of the null distribution, and remains an open problem. It appears that MMDu and MMDl fall at two extremes of a spectrum: the former has the lowest variance of any n -sample estimator, and should be used in limited data regimes; the latter is the estimator requiring the least computation while still looking at each of the samples, and usually achieves better Type II error than MMDu at a given computational cost, albeit by looking at much more data (the ‘‘limited time, unlimited data’’ scenario). A major reason MMDl is faster is that its null distribution is straightforward to compute, since it is Gaussian and its variance can be calculated at the same cost as the test statistic. A reasonable next step would be to find a compromise between these two extremes: to construct a statistic with a lower variance than MMDl, while retaining an asymptotically Gaussian null distribution (hence remaining faster than tests based on MMDu). We study a family of such test statistics, where we split the data into blocks of size B , compute the quadratic-time MMDu on each block, and then average the resulting statistics. We call the resulting tests

B-tests. As long as we choose the size B of blocks such that $n/B \rightarrow \infty$, we are still guaranteed asymptotic Normality by the central limit theorem, and the null distribution can be computed at the same cost as the test statistic. For a given sample size n , however, the power of the test can increase dramatically over the MMD test, even for moderate block sizes B , making much better use of the available data with only a small increase in computation. The block averaging scheme was originally proposed in [13], as an instance of a two-stage U-statistic, to be applied when the degree of degeneracy of the U-statistic is indeterminate. Differences with respect to our method are that Ho and Shieh compute the block statistics by sampling with replacement [13, (b) p. 863], and propose to obtain the variance of the test statistic via Monte Carlo, jackknife, or bootstrap techniques, whereas we use closed form expressions. Ho and Shieh further suggest an alternative two-stage U-statistic in the event that the degree of degeneracy is known; we return to this point in the discussion. While we confine ourselves to the MMD in this paper, we emphasize that the block approach applies to a much broader variety of test situations where the null distribution cannot easily be computed, including the energy distance and distance covariance [18, 2, 22] and Fisher statistic [12] in the case of two-sample testing, and the Hilbert-Schmidt Independence Criterion [8] and distance covariance [23] for independence testing. Finally, the kernel learning approach of [9] applies straightforwardly, allowing us to maximize test power over a given kernel family. Code is available at <http://github.com/wojzaremba/btest>.

2

Theory

In this section we describe the mathematical foundations of the B-test. We begin with a brief review of kernel methods, and of the maximum mean discrepancy. We then present our block-based average MMD statistic, and derive its distribution under the H_0 ($P = Q$) and H_A ($P \neq Q$) hypotheses. The central idea employed in the construction of the B-test is to generate a low variance MMD estimate by averaging multiple low variance kernel statistics computed over blocks of samples. We show simple sufficient conditions on the block size for consistency of the estimator. Furthermore, we analyze the properties of the finite sample estimate, and propose a consistent strategy for setting the block size as a function of the number of samples. 2.1

Definition and asymptotics of the block-MMD

Let F_k be an RKHS defined on a topological space X with reproducing kernel k , and P a Borel probability measure on X . The mean embedding of P in F_k , written $\mu_k(P)$, is defined such that

250 H_A histogram

250

H_0 histogram

H_A histogram

approximated 5% quantile of H_0

H histogram 0

approximated 5% quantile of H

200

0
 200
 150 150
 100 100
 50 50
 0 ?4 0 ?0.05
 ?0.04
 ?0.03
 ?0.02
 ?0.01
 0
 0.01
 0.02
 0.03
 0.04
 ?2
 0
 2
 4
 6
 8
 10 ?3
 x 10
 0.05

- (a) $B = 2$. This setting corresponds to the MMDl statistic [10].
 (b) $B = 250$

Figure 1: Empirical distributions under H_0 and H_A for different regimes of B for the music experiment (Section 3.2). In both plots, the number of samples is fixed at 500. As we vary B , we trade off the quality of the finite sample Gaussian approximation to the null distribution, as in Theorem 2.3, with the variances of the H_0 and H_A distributions, as outlined in Section 2.1. In (b) the distribution under H_0 does not resemble a Gaussian (it does not pass a level 0.05 Kolmogorov-Smirnov (KS) normality test [16, 20]), and a Gaussian approximation results in a conservative test threshold (vertical green line). The remaining empirical distributions all pass a KS normality test.

that $\mathbb{E}x^p f(x) = hf, \sum_k (p)_i F_k$ for all $f \in F_k$, and exists for all Borel probability measures when k is bounded and continuous [3, 10]. The maximum mean discrepancy (MMD) between a Borel probability measure P and a second Borel probability measure Q is the squared RKHS distance between their respective mean embeddings, 2

$$\sum_k (P, Q) = k \sum_k (P) \sum_k (Q) k F_k = \mathbb{E}x_0 k(x, x_0) + \mathbb{E}y_0 k(y, y_0) - 2\mathbb{E}x_0 y_0 k(x, y),$$

(1)
0

where x denotes an independent copy of x [11]. Introducing the notation $z = (x, y)$, we write $\sum_k (P, Q) = \mathbb{E}z_0 h_k(z, z_0)$,

$$h(z, z_0) = k(x, x_0) + k(y, y_0) - k(x, y_0) - k(x_0, y). \quad (2)$$

When the kernel k is characteristic, then $\gamma_k(P, Q) = 0$ iff $P = Q$ [21]. Clearly, the minimum variance unbiased estimate MMD_u of $\gamma_k(P, Q)$ is a U-statistic. By analogy with MMD_u , we make use of averages of $h(x, y, x_0, y_0)$ to construct our two-sample test. We denote by $\gamma_k^{(i)}$ the i th empirical estimate MMD_u based on a subsample of size B , where $n \geq B$ (for notational purposes, we will index samples as though they are presented in a random fixed order). More precisely, $\gamma_k^{(i)} =$

$$\frac{1}{B(B-1)} \sum_{i \neq j} h(z_a, z_b). \quad (3)$$

$$a = (i-1)B+1, b = (i-1)B+1+B = a+B$$

The B-test statistic is an MMD estimate obtained by averaging the $\gamma_k^{(i)}$. Each $\gamma_k^{(i)}$ under H_0 converges to an infinite sum of weighted χ^2 variables [7]. Although setting $B = n$ would lead to the lowest variance estimate of the MMD, computing sound thresholds for a given p-value is expensive, involving repeated bootstrap sampling [5, 14], or computing the eigenvalues of a Gram matrix [7]. In contrast, we note that $\gamma_k^{(i)}$, $i=1, \dots, n$ are i.i.d. variables, and averaging them allows us to apply the central limit theorem in order to estimate p-values from a normal distribution. We denote the average of the $\gamma_k^{(i)}$ by $\bar{\gamma}_k$, $\bar{\gamma}_k = \frac{1}{n} \sum_{i=1}^n \gamma_k^{(i)}$. (4) We would like to apply the central limit theorem to variables $\gamma_k^{(i)}$, $i=1, \dots, n$. It remains for us to derive the distribution of γ_k under H_0 and under H_A . We rely on the result from [11, Theorem 8] for H_A . According to our notation, for every i ,

Theorem 2.1 Assume $0 \leq E(h^2) < \infty$, then under H_A , γ_k converges in distribution to a Gaussian according to

$$\sqrt{B}(\bar{\gamma}_k - MMD^2) \xrightarrow{D} N(0, \sigma^2), \quad \text{where } \sigma^2 = 4 E_{z, z_0} [(E_{z_0} h(z, z_0))^2 - E_{z, z_0} (h(z, z_0))^2]. \quad (5)$$

This in turn implies that

$\bar{\gamma}_k \xrightarrow{D} N(MMD^2, \sigma^2/B)$. For an average of $\{\gamma_k^{(i)}\}_{i=1, \dots, Bn}$, the central limit theorem implies that under H_A ,

$$\sqrt{Bn}(\bar{\gamma}_k - MMD^2) \xrightarrow{D} N(MMD^2, \sigma^2/Bn). \quad (6)$$

$$(7)$$

This result shows that the distribution of H_A is asymptotically independent of the block size, B . Turning to the null hypothesis, [11, Theorem 8] additionally implies that under H_0 for every i , Theorem 2.2

$$\sqrt{B} \gamma_k^{(i)} \xrightarrow{D} N(0, \sigma^2), \quad \text{where } \sigma^2 \leq 2, \quad (8)$$

bias caused the tests to be overly conservative, with lower Type I error than the design level required (Figures 2 and 5). 2.3

Finite Sample Case

In the finite sample case, we apply the Berry-Esseen theorem, which gives conservative bounds on the ϵ convergence of a series of finite sample random variables to a Gaussian distribution [4]. 2.2 Theorem 2.3 Let X_1, X_2, \dots, X_n be i.i.d. variables. $E(X) \leq 0$, and $E(X^2) = 1$, $E(X^3) = \gamma$. Let F_n be a cumulative distribution of $\sum_{i=1}^n X_i$, and let Φ denote the standard normal distribution. Then for every x ,

$$|F_n(x) - \Phi(x)| \leq C \frac{|\gamma|}{\sqrt{n}} \quad (12)$$

where $C \leq 1$. This result allows us to ensure fast point-wise convergence of the B-test. We have that $\epsilon(\epsilon, k) = O(1)$, i.e., it is dependent only on the underlying distributions of the samples and not on the sample size. The number of i.i.d. samples is $nB \geq 1$. Based on Theorem 2.3, the point-wise error can be $O(1/B^2)$ upper bounded by $O(\epsilon)$ under H_A . Under H_0 , the error can be bounded by $O(\epsilon/n) = O(1/nB^2)$.

$$O(1/nB^2) = O(1/nB^2)$$

$$\epsilon_n =$$

$$)2$$

$$B$$

$$3.5 O(1/B^2 \epsilon_n).$$

$$B$$

While the asymptotic results indicate that convergence to an optimal predictor is fastest for larger B , the finite sample results support decreasing the size of B in order to have a sufficient number n of samples for application of the central limit theorem. As long as $B \geq 1$ and $B \geq 1/\epsilon$, the assumptions of the B-test are fulfilled. By varying B , we make a fundamental tradeoff in the construction of our two sample test. When B is small, we have many samples, hence the null distribution is close to the asymptotic limit provided by the central limit theorem, and the Type I error is estimated accurately. The disadvantage of a small B is a lower test power for a given sample size. Conversely, if we increase B , we will have a lower variance empirical distribution for H_0 , hence higher test power, but we may have a poor estimate of the number of Type I errors (Figure 1). A sensible family of heuristics therefore is to set $B = \lceil n^\epsilon \rceil$

$$(13)$$

for some $0 \leq \epsilon \leq 1$, where we round to the nearest integer. In this setting the number of samples (nB) available for application of the central limit theorem will be $\lceil n^{1+\epsilon} \rceil$. For given ϵ , the complexity of the B-test is $O(n^{1+\epsilon})$. We note that any value of $\epsilon \in (0, 1)$ yields a consistent estimator. We have chosen $\epsilon = 2$ in the experimental results section, with resulting complexity $O(n^{1.5})$: we emphasize that this is a heuristic, and just one choice that fulfils our assumptions.

$$3$$

Experiments

We have conducted experiments on challenging synthetic and real datasets in order to empirically measure (i) sample complexity, (ii) computation time, and (iii) Type I / Type II errors. We evaluate B-test performance in comparison to the MMDl and MMDu estimators, where for the latter we compare across different strategies for null distribution quantile estimation. 5

Method									
Kernel parameters	?=1								
B-test									
?	= median	multiple	kernels						
Pearson curves	Gamma approximation	Gram matrix	spectrum	Bootstrap					
Pearson curves	Gamma approximation	Gram matrix	spectrum	Bootstrap					
Additional parameters	B=2	B=?8	B= n	any B	B=2	B =p8	B = n2		
Minimum number of samples	26400	3850	886	¿	60000	37000	5400	1700	
Computation time (s)	0.0012	0.0039	0.0572						
B=n									
	186	183	186	190					
	387.4649	0.2667	407.3447	129.4094					
?=1									
Consistent	X	X	X	X	X	X	X	X	
	0.0700	0.1295	0.8332						
?	?	X	X	?	?	X	X		
¿	60000,	or 2h	per iteration	timeout					
?	= median								

Table 1: Sample complexity for tests on the distributions described in Figure 3. The fourth column indicates the minimum number of samples necessary to achieve Type I and Type II errors of 5%. The fifth column is the computation time required for 2000 samples, and is not presented for settings that have unsatisfactory sample complexity. 0.08

0.08									
Empirical Type I error	Expected	?	Type I error	B=	n				
0.07									
0.08									
Empirical Type I error	Expected	?	Type I error	B=	n				
0.07									
0.06									
0.06									
0.05									
0.05									
0.05									
0.04	0.03								
0.04	0.03								
0.02									
0.02									
0.01									
0.01									
0	2								

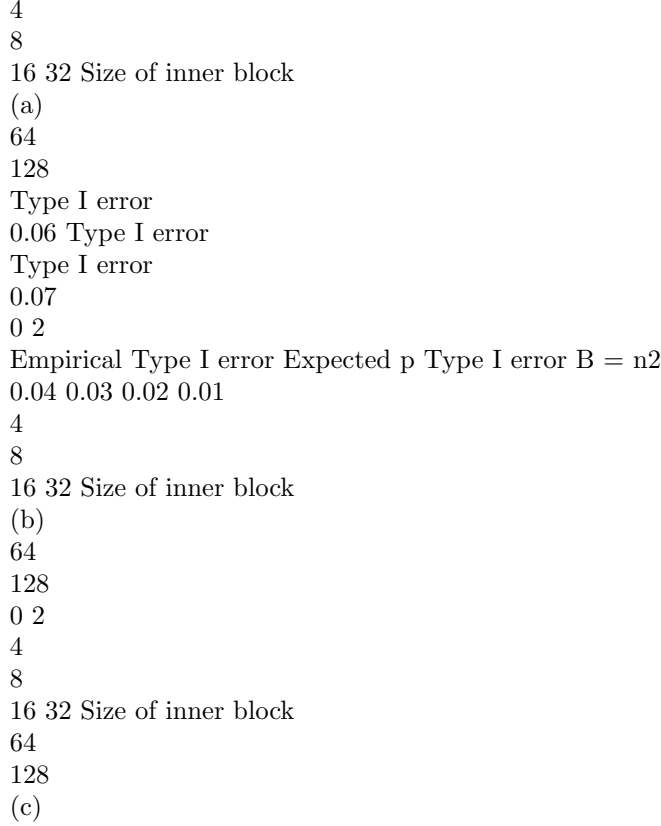


Figure 2: Type I errors on the distributions shown in Figure 3 for $\alpha = 5\%$: (a) MMD, single kernel, $\alpha = 1$, (b) MMD, single kernel, α set to the median pairwise distance, and (c) MMD, non-negative linear combination of multiple kernels. The experiment was repeated 30000 times. Error bars are not visible at this scale.

3.1

Synthetic data

Following previous work on kernel hypothesis testing [9], our synthetic distributions are 5×5 grids of 2D Gaussians. We specify two distributions, P and Q . For distribution P each Gaussian has identity covariance matrix, while for distribution Q the covariance is non-spherical. Samples drawn from P and Q are presented in Figure 3. These distributions have proved to be very challenging for existing non-parametric two-sample tests [9]. We employed three different kernel selection strategies in the hypothesis test. First, we used a Gaussian kernel with $\alpha = 1$, which approximately matches the scale of the variance of each Gaussian in mixture P . While this is a somewhat arbitrary default choice, we selected it as it performs well in practice (given the lengthscale of the (a) Distribution P (b) Distribution Q data), and we treat it as a baseline. Next, we set α equal to the median pairwise distance over the training data, Figure 3: Synthetic data distributions P and Q which is a standard way to choose the Gaussian

kernel Q . Samples belonging to these classes are bandwidth [17], although it is likewise arbitrary in this difficult to distinguish. context. Finally, we applied a kernel learning strategy, in which the kernel was optimized to maximize the test power for the alternative $P \neq Q$ [9]. This approach returned a non-negative linear combination combination of base kernels, where half the data were used in learning the kernel weights (these data were excluded from the testing phase). The base kernels in our experiments were chosen to be Gaussian, with bandwidths in the set $\{2^{15}, 2^{14}, \dots, 2^{10}\}$. Testing was conducted using the remaining half of the data. 6

B -test, a single kernel, $\gamma = 1$ B -test, a single kernel, $\gamma = \text{median } B$ -test
kernel selection Tests estimating MMDu with $\gamma=1$

Tests estimating MMDu with $\gamma=\text{median}$ Empirical number of Type II errors

For comparison with the quadratic time U statistic MMDu [7, 10], we evaluated four null distribution estimates: (i) Pearson curves, (ii) gamma approximation, (iii) Gram matrix spectrum, and (iv) bootstrap. For methods using Pearson curves and the Gram matrix spectrum, we drew 500 samples from the null distribution estimates to obtain the 1 - α quantiles, for a test of level α . For the bootstrap, we fixed the number of shuffles to 1000. We note that Pearson curves and the gamma approximation are not statistically consistent. We considered only the setting with $\gamma = 1$ and γ set to the median pairwise distance, as kernel selection is not yet solved for tests using MMDu [9].

1 0.8 0.6 0.4 0.2 0

1

10

2

10 Size of inner block

3

10

Figure 4: Synthetic experiment: number of Type II er-

In the first experiment we set the Type I error to α vs B , given a fixed probability γ of Type I error 5%, and we recorded the Type II error. We α . As B grows, the Type II error drops quickly when conducted these experiments on 2000 samples the kernel is appropriately chosen. The kernel selection over 1000 repetitions, with varying block size, selection method is described in [9], and closely approximates the baseline performance of the well-informed choice strategies, as a function of B . The median heuristic performs extremely poorly in this experiment. As discussed in [9, Section 5], the reason for this failure is that the lengthscale of the difference between the distributions P and Q differs from the lengthscale of the main data variation as captured by the median, which gives too broad a kernel for the data. In the second experiment, our aim was to compare the empirical sample complexity of the various methods. We again fixed the same Type I error for all methods, but this time we also fixed a Type II error of 5%, increasing the number of samples until the latter error rate was achieved. Column four of Table 1 shows the number of samples required in each setting to achieve these error rates. We additionally compared the computational efficiency of the

various methods. The computation time for each method with a fixed sample size of 2000 is presented in column five of Table 1. All experiments were run on a single 2.4 GHz core. Finally, we evaluated the empirical Type I error for $\alpha = 5\%$ and increasing B . Figure 2 displays the empirical Type I error, where we note the location of the $\alpha = 0.5$ heuristic in Equation (13). For the user-chosen kernel ($\alpha = 1$, Figure 2(a)), the number of Type I errors closely matches the targeted test level. When median heuristic is used, however, the test is overly conservative, and makes fewer Type I errors than required (Figure 2(b)). This indicates that for this choice of α , we are not in the asymptotic regime, and our Gaussian null distribution approximation is inaccurate. Kernel selection via the strategy of [9] alleviates this problem (Figure 2(c)). This setting coincides with a block size substantially larger than 2 (MMD1), and therefore achieves lower Type II errors while retaining the targeted Type I error. 3.2

Musical experiments

In this set of experiments, two amplitude modulated Rammstein songs were compared (Sehnsucht vs. Engel, from the album Sehnsucht). Following the experimental setting in [9, Section 5], samples from P and Q were extracts from AM signals of time duration 8.3×10^{-3} seconds in the original audio. Feature extraction was identical to [9], except that the amplitude scaling parameter α was set to 0.3 instead of 0.5. As the feature vector had size 1000 we set the block size $B = 1000 = 32$. Table 2 summarizes the empirical Type I and Type II errors over 1000 repetitions, and the average computation times. Figure 5 shows the average number of Type I errors as a function of B : in this case, all kernel selection strategies result in conservative tests (lower Type I error than required), indicating that more samples are needed to reach the asymptotic regime. Figure 1 shows the empirical H_0 and H_A distributions for different B .

4

Discussion

We have presented experimental results both on a difficult synthetic problem, and on real-world data from amplitude modulated audio recordings. The results show that the B-test has a much better 7

Kernel parameters

Method

Additional parameters $B = 2$ $B = n$ $B = 2n$ $B = n^2$ $B = n^2$

$\alpha = 1$ B-test

$\alpha = \text{median}$ multiple kernels

Gram matrix spectrum Bootstrap Gram matrix spectrum Bootstrap

Type I error

Type II error

0.038 0.006 0.043 0.026 0.0481 0.025

0.927 0.597 0.786 0 0.867 0.012

Computational time (s) 0.039 1.276 0.047 1.259 0.607 18.285

0 0.01 0 0.01

0 0 0 0

160.1356 121.2570 286.8649 122.8297

$\alpha = 1$ $B = 2000$ $\alpha = \text{median}$

Table 2: A comparison of consistent tests on the music experiment described in Section 3.2. Here computation time is reported for the test achieving the stated error rates.

0.08
0.08
Empirical Type I error Expected ? Type I error B= n
0.07
0.08
Empirical Type I error Expected ? Type I error B= n
0.07
0.06
0.06
0.05
0.05
0.05
0.04 0.03
Type I error
0.06 Type I error
Type I error
0.07
0.04 0.03
0.04 0.03
0.02
0.02
0.02
0.01
0.01
0.01
0 2
4
8
16 32 Size of inner block
(a)
64
128
0 2
4
8
16 32 Size of inner block
(b)
64
128
Empirical Type I error Expected p Type I error B = n2
0 2
4
8

16 32 Size of inner block

64

128

(c)

Figure 5: Empirical Type I error rate for $\alpha = 5\%$ on the music data (Section 3.2). (a) A single kernel test with $\alpha = 1$, (b) A single kernel test with $\alpha = \text{median}$, and (c) for multiple kernels. Error bars are not visible at this scale. The results broadly follow the trend visible from the synthetic experiments.

sample complexity than MMDI over all tested kernel selection strategies. Moreover, it is an order of magnitude faster than any test that consistently estimates the null distribution for MMDu (i.e., the Gram matrix eigenspectrum and bootstrap estimates): these estimates are impractical at large sample sizes, due to their computational complexity. Additionally, the B-test remains statistically consistent, with the best convergence rates achieved for large B. The B-test combines the best features of MMDI and MMDu based two-sample tests: consistency, high statistical efficiency, and high computational efficiency. A number of further interesting experimental trends may be seen in these results. First, we have observed that the empirical Type I error rate is often conservative, and is less than the 5% targeted by the threshold based on a Gaussian null distribution assumption (Figures 2 and 5). In spite of this conservatism, the Type II performance remains strong (Tables 1 and 2), as the gains in statistical power of the B-tests improve the testing performance (cf. Figure 1). Equation (7) implies that the size of B does not influence the asymptotic variance under H_A , however we observe in Figure 1 that the empirical variance of H_A drops with larger B. This is because, for these P and Q and small B, the null and alternative distributions have considerable overlap. Hence, given the distributions are effectively indistinguishable at these sample sizes n, the variance of the alternative distribution as a function of B behaves more like that of H_0 (cf. Equation (10)). This effect will vanish as n grows. Finally, [13] propose an alternative approach for U-statistic based testing when the degree of degeneracy is known: a new U-statistic (the TU-statistic) is written in terms of products of centred U-statistics computed on the individual blocks, and a test is formulated using this TU-statistic. Ho and Shieh show that a TU-statistic based test can be asymptotically more powerful than a test using a single U-statistic on the whole sample, when the latter is degenerate under H_0 , and nondegenerate under H_A . It is of interest to apply this technique to MMD-based two-sample testing. Acknowledgments We thank Mladen Kolar for helpful discussions. This work is partially funded by ERC Grant 259112, and by the Royal Academy of Engineering through the Newton Alumni Scheme.

8

2 References

- [1] Bengt Von Bahr. On the convergence of moments in the central limit theorem. The Annals of Mathematical Statistics, 36(3):pp. 808?818, 1965. [2] L.

Baringhaus and C. Franz. On a new multivariate two-sample test. *J. Multivariate Anal.*, 88:190?206, 2004. [3] A. Berlinet and C. Thomas-Agnan. *Reproducing Kernel Hilbert Spaces in Probability and Statistics*. Kluwer, 2004. [4] Andrew C Berry. The accuracy of the gaussian approximation to the sum of independent variates. *Transactions of the American Mathematical Society*, 49(1):122?136, 1941. [5] B. Efron and R. J. Tibshirani. *An Introduction to the Bootstrap*. Chapman & Hall, 1993. [6] M. Fromont, B. Laurent, M. Lerasle, and P. Reynaud-Bouret. Kernels based tests with nonasymptotic bootstrap approaches for two-sample problems. In *COLT*, 2012. [7] A. Gretton, K. Fukumizu, Z. Harchaoui, and B. Sriperumbudur. A fast, consistent kernel twosample test. In *Advances in Neural Information Processing Systems 22*, pages 673?681, 2009. [8] A. Gretton, K. Fukumizu, C.-H. Teo, L. Song, B. Sch?olkopf, and A. J. Smola. A kernel statistical test of independence. In *Advances in Neural Information Processing Systems 20*, pages 585?592, Cambridge, MA, 2008. MIT Press. [9] A. Gretton, B. Sriperumbudur, D. Sejdinovic, H. Strathmann, S. Balakrishnan, M. Pontil, and K. Fukumizu. Optimal kernel choice for large-scale two-sample tests. In *Advances in Neural Information Processing Systems 25*, pages 1214?1222, 2012. [10] Arthur Gretton, Karsten M. Borgwardt, Malte J. Rasch, Bernhard Sch?olkopf, and Alexander Smola. A kernel two-sample test. *J. Mach. Learn. Res.*, 13:723?773, March 2012. [11] Arthur Gretton, Karsten M. Borgwardt, Malte J. Rasch, Bernhard Sch?olkopf, and Alexander J. Smola. A kernel method for the two-sample-problem. In *NIPS*, pages 513?520, 2006. [12] Z. Harchaoui, F. Bach, and E. Moulines. Testing for homogeneity with kernel Fisher discriminant analysis. In *NIPS*, pages 609?616. MIT Press, Cambridge, MA, 2008. [13] H.-C. Ho and G. Shieh. Two-stage U-statistics for hypothesis testing. *Scandinavian Journal of Statistics*, 33(4):861?873, 2006. [14] Norman Lloyd Johnson, Samuel Kotz, and Narayanaswamy Balakrishnan. *Continuous univariate distributions*. *Distributions in statistics*. Wiley, 2nd edition, 1994. [15] A. Kleiner, A. Talwalkar, P. Sarkar, and M. I. Jordan. A scalable bootstrap for massive data. *Journal of the Royal Statistical Society, Series B*, In Press. [16] Andrey N Kolmogorov. Sulla determinazione empirica di una legge di distribuzione. *Giornale dell'Istituto Italiano degli Attuari*, 4(1):83?91, 1933. [17] B. Sch?olkopf. *Support vector learning*. Oldenbourg, M?unchen, Germany, 1997. [18] D. Sejdinovic, A. Gretton, B. Sriperumbudur, and K. Fukumizu. Hypothesis testing using pairwise distances and associated kernels. In *ICML*, 2012. [19] R. Serfling. *Approximation Theorems of Mathematical Statistics*. Wiley, New York, 1980. [20] Nickolay Smirnov. Table for estimating the goodness of fit of empirical distributions. *The Annals of Mathematical Statistics*, 19(2):279?281, 1948. [21] B. Sriperumbudur, A. Gretton, K. Fukumizu, G. Lanckriet, and B. Sch?olkopf. Hilbert space embeddings and metrics on probability measures. *Journal of Machine Learning Research*, 11:1517?1561, 2010. [22] G. Sz?ekely and M. Rizzo. Testing for equal distributions in high dimension. *InterStat*, (5), November 2004. [23] G. Sz?ekely, M. Rizzo, and N. Bakirov. Measuring and testing dependence by correlation of distances. *Ann. Stat.*, 35(6):2769?2794, 2007. [24] M. Yamada, T. Suzuki, T. Kanamori, H. Hachiya, and M. Sugiyama. Relative density-ratio estimation for robust distribution comparison. *Neural*

Computation, 25(5):1324?1370, 2013.
9